Processing and Visualising Large DEMs on a Laptop

Progress Report, Phase 5: Interactive Viewing with Debounced Pan/Zoom and Prefetch

Team GeoGuessers | Spatial Informatics

Swayam Goyal (2024111012) Siddharth Choudhary (2024111021)
Yakkala Viswanadha Kartikeya (2024111033)

Abstract

Phases 2–4 showed that terrain processing on a laptop is bound by decoding and memory, not by arithmetic, and that a derivative computed on demand from a pyramid level costs the same few milliseconds at every zoom level. This phase builds a viewer on that result and measures it. A tile server computes slope, aspect, hillshade and roughness products for fixed 512×512 tiles on demand. Its foreground workers run on the performance cores and prefetch on the efficiency cores, with block and tile caches and cancellation. A browser viewer pans and zooms continuously: it commits a zoom once, after a trailing debounce, and prefetches the eight neighbouring views. A QGIS plugin computes any dragged box. We replayed 17 seeded interaction traces against a fresh server in 24 configurations, 1,068 runs in all. With prefetch, 98–100% of the tiles entering the view while panning are already in the browser, at up to two view widths per second. The debounce cuts the tiles the server renders per zoom gesture from 62 to 19, and the first exact tile follows the debounce within 63–71 ms. Prefetch does not slow the foreground: on identical tiles the server's 95th-percentile latency is 90 ms with prefetch and 95 ms without, because prefetch stops whenever the user acts. Notched-wheel traces set the debounce default to 150 ms. A real browser (headless Chrome) reproduces the simulated client's request counts and timings within 10%.

1 Objective

The meeting asked for a viewer that works like a map: pan and zoom freely, compute only what is in view at the resolution the screen needs, and prefetch what the user will probably look at next. The goals come from the team's Phase 5 specification; the internal design is ours. The questions are:

1. Can on-the-fly computation keep up with continuous pan and zoom on one laptop?
2. Does a trailing debounce on zoom cut the server's work, and what delay does it add?
3. Does prefetching the neighbouring views make panning seamless without slowing the visible tiles?
4. Which settings (debounce, prefetch policy, worker pools, core placement, caches) should be the defaults?

Two gates were set in advance: prefetch must not raise foreground p95 latency by more than 10%, and the debounce must reduce server tiles per zoom gesture while keeping the first frame within 100 ms of the debounce.

2 The system

Table 1: Components built in this phase (all in the repository).

Component	Design
Tile engine	Fixed 512×512 tile lattice per pyramid level, aligned to the COG's 512^2 blocks, with a 1-cell halo. One tile is rendered on one thread by GIL-free Numba kernels, so tiles run in parallel (5.4× on 8 threads). Decoded blocks are cached and shared by neighbouring tiles; encoded tiles are cached too. Output is 8-bit palette PNG. Modes: *auto* (the level's own cells, E5 error), *exact* (full resolution, averaged), *precomputed*.
Scheduler	Foreground and background queues. Foreground workers are pinned to the 12 performance-core threads and background workers to the 8 efficiency cores. Identical requests share one job; a foreground request promotes a queued prefetch job. A newer generation of a session supersedes older work. Cancellation is checked between read, kernel and encode. Prefetch does not start while foreground work is pending, and pauses when memory is low.
Server	FastAPI: tiles, overview, `POST /view` (a box as PNG or GeoTIFF), session generations, stats. Every tile reports its queue, read, decode, kernel, colour-map and encode times. A closed connection cancels its job (median 37 ms).
Web viewer	Each frame draws the needed level's tiles, falling back to a coarser or finer cached level and then the overview. Zoom commits once, after a trailing debounce; pan requests new tiles at once. After a commit the eight neighbouring views are prefetched, in the direction of travel first, at most two at a time. Shift + drag zooms to a box.
QGIS plugin	Drag a box; the server's GeoTIFF result is added as a styled layer. A full-resolution box is identical to `gdaldem`'s slope of the full-resolution DEM.

3 Method

Virtual client. A Python port of the viewer's policy (camera, level choice, browser cache, debounce, generations, aborts, prefetch) reads the same policy file as the browser. It replays a trace in real time against the live server, over HTTP/1.1 with at most six connections (a browser's limit per server), and simulates 60 Hz frames. Python's default event-loop clock on Windows advances in 15.6 ms steps, so the client runs its own loop on the high-resolution clock.

Traces. Seeded, for a 1600×900 view:

- **Pans:** straight and reversing pans at 0.5, 1 and 2 view widths per second (5 s), and a random walk.
- **Zoom bursts:** 60 Hz wheel bursts of 0.3, 1 and 2 s, in and out, over 1–3 levels.
- **Notched wheel:** ten discrete notches 30–150 ms or 120–300 ms apart.
- **Zoom-pan:** 10 s of alternating bursts and pans.
- **Box:** three nested box zooms from the whole map to full resolution.

Protocol. One factor at a time around the defaults, 24 configurations; 1 warm-up and 5 measured runs for every configuration and trace. Each run:

- starts a new server process with empty caches, and loads the first view and its prefetch ring before the trace clock starts;
- reads the server's counters at the start and end of the trace, and samples memory every 50 ms.

The data is the Chamoli 2 m DEM (pyramid levels 0–6), with hillshade, unless stated.

Metrics.

- *Pan hit rate*: share of tiles already cached when they enter the view.
- *Blank area-time*: share of view area × frames not drawn from the needed level's own tiles.
- *TTFF*: last wheel tick to the first frame showing a tile of the committed level.
- *Complete*: last tick to the first frame in which every visible tile is drawn.
- *Server tiles per gesture*: foreground jobs finished or cancelled mid-render.
- *Prefetch precision*: prefetched tiles later drawn, as a share of all prefetched tiles.
- *Foreground p95*: measured at the client and on the server.

4 Results

4.1 Defaults

Table 2: Default policy on every trace (median of 5 runs; debounce 200 ms, the sweep's base). TTFF and Complete (ms) for zoom traces include the debounce; Tiles/gest. = server foreground tiles per zoom gesture.

Trace	Fg req.	Prefetched	Hit rate	TTFF	Complete	Tiles/gest.
pan-line 0.5/1/2	0 / 0 / 2	175 / 313 / 333	1.00 / 1.00 / 0.99			
pan-reverse 0.5/1/2	0 / 0 / 5	73 / 128 / 443	1.00 / 1.00 / 0.98			
pan-walk	0	290	1.00			
zoom in 0.3/1/2 s	20 / 15 / 20			272 / 273 / 289	325 / 325 / 358	20 / 15 / 20
zoom out 0.3/1/2 s	18 / 20 / 12			269 / 264 / 257	321 / 314 / 276	18 / 20 / 12
zoom-pan	36	128	1.00	206	206	7.2
box	35	125		64	127	

On-the-fly computation keeps up (Table 2):

- **Pans** need almost no foreground requests: the prefetch ring already holds the tiles.
- **Zooms:** the first exact tile arrives 57–89 ms after the commit and the full view 76–158 ms after it.
- **Box zoom** shows its first exact tile 64 ms after the animation ends.
- **Cost:** the server never needed more than 982 MiB. Prefetch used about 3 CPU-seconds of the efficiency cores per 5 s pan.

Figure 1 shows what the viewer draws during zoom-pan: the needed level's own tiles cover the view except for a few hundred milliseconds after each zoom-in, when cached coarser tiles stand in.

4.2 Debounce (gate 2)

Table 3: Server foreground tiles per zoom gesture and delay, by debounce (median of runs).

Debounce	60 Hz bursts	Notched 30–150 ms	Notched 120–300 ms	TTFF − debounce, bursts (ms)	TTFF, bursts (ms)
off (every tick)	62	36	48	–	6
100 ms	19	36	48	63	163
150 ms (new default)	19	**15**	**39**	68	218
200 ms	19	15	39	71	271
250 ms	19	15	39	68	318
200 ms + forced commit at 500 ms	35	30	39	62	262

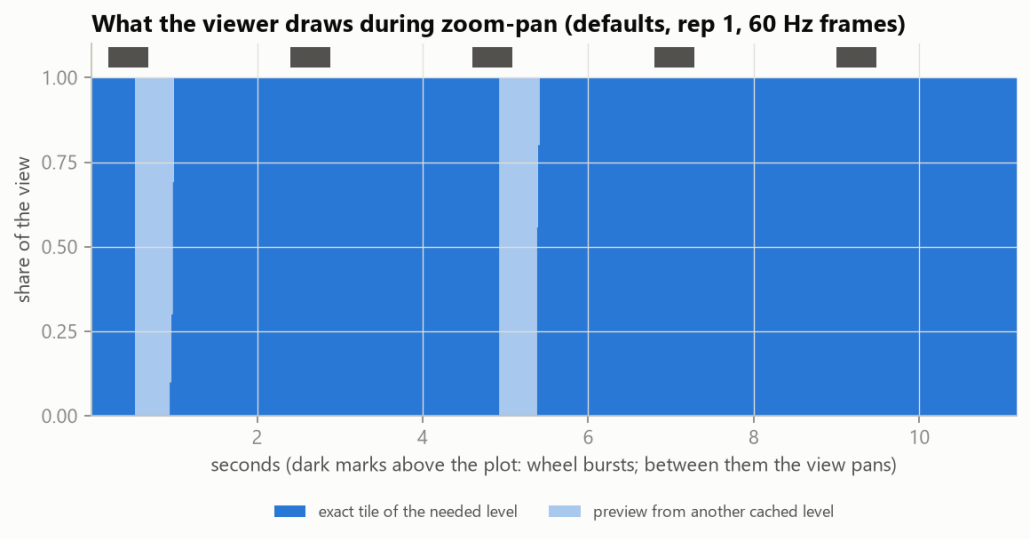


Figure 1: Share of the view drawn from the needed level's own tiles (dark) or from another cached level (light), 60 Hz frames, during zoom-pan with the defaults.

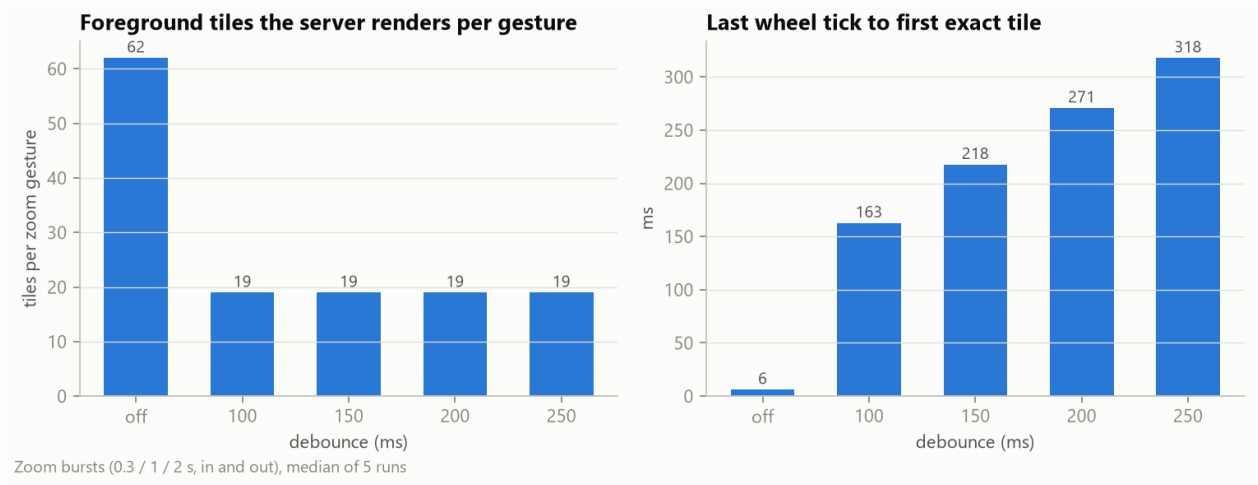


Figure 2: Zoom bursts at 60 Hz: foreground tiles the server renders per gesture, and time from the last tick to the first exact tile, against the debounce.

- **Without a debounce** every tick commits, up to 121 times in a 2 s burst. The server renders a median of 62 foreground tiles per gesture and cancels 4–19 of them mid-render.
- **With a debounce** it renders 19, **3.3× fewer**, and the first exact tile arrives 63–71 ms after the debounce fires. The gate passes (Table 3, Figure 2).
- **Why notched wheels decide the value.** At 60 Hz the ticks are 17 ms apart, so every debounce from 100 to 250 ms behaves the same. A notched mouse wheel, with gaps of 30–300 ms, separates them:
  - 100 ms commits between notches and does as badly as no debounce (36 tiles);
  - 150 ms and longer coalesce the gesture (15 tiles);
  - 200 ms only adds 40–50 ms of delay over 150 ms.

  The default is now 150 ms.
- **Forced commit:** committing after 500 ms of continuous scrolling almost doubles the tiles for 9 ms less delay, so it stays off.

4.3 Prefetch

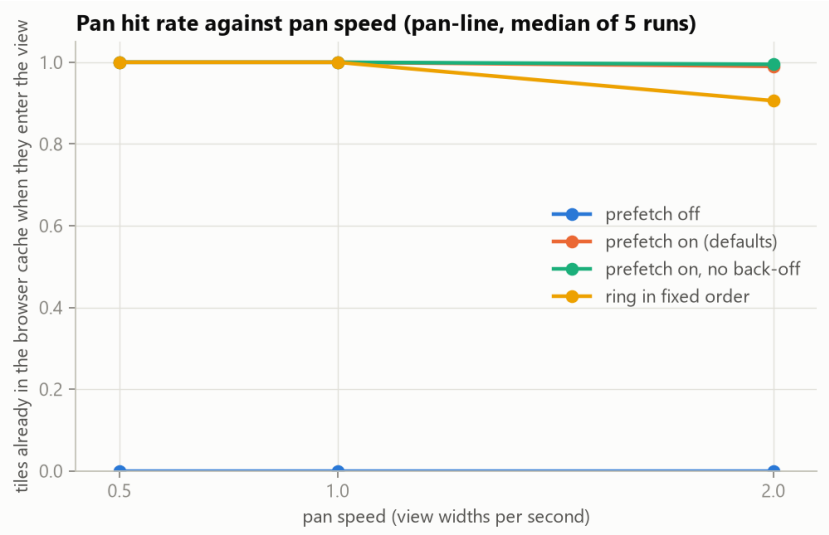


Figure 3: Share of tiles already in the browser when they enter the view, straight pans, by pan speed.

Table 4: Prefetch policies over the seven pan traces and zoom-pan (40 runs each; median [interquartile range]).

Policy	Foreground requests	Pan hit rate	Prefetched	Precision
prefetch off	104 [68–135]	0.30 [0.00–0.51]	0	–
defaults (back-off if foreground p95 > 250 ms)	0 [0–4]	1.00 [1.00–1.00]	218	0.34
first policy (back-off at 100 ms)	0 [0–56]	1.00 [0.70–1.00]	192	0.29
no back-off	0 [0–4]	1.00 [1.00–1.00]	232	0.34
ring in fixed order	0 [0–7]	1.00 [0.97–1.00]	292	0.33

- **With prefetch** the tiles entering the view during a pan are 98–100% cached at every speed. The pans need 0–5 foreground requests instead of 36–202 (Table 4, Figure 3).

- **Without prefetch** the hit rate is 0% on a straight pan and about 50% when the pan turns back. The view still looks almost complete (blank area-time 0.2–3.4%), because each new tile arrives in about 50 ms and a coarser level covers the gap.
- **Cost:** the ring fetches 2–5 tiles for each one used (precision 0.20–0.56), since it always loads directions the user does not take.
- **Ordering:** loading the direction of travel first keeps the hit rate at 0.99 at two view widths per second; a fixed order drops it to 0.91.
- **Back-off threshold:** the first policy stopped prefetch for 2 s whenever foreground p95 exceeded 100 ms. That is the cold-start latency, so prefetch switched off at random, and the fast reversing pan fell to a 0.50 hit rate. At 250 ms the policy matches "no back-off" everywhere and keeps a safety valve.

4.4 Prefetch and the foreground (gate 1)

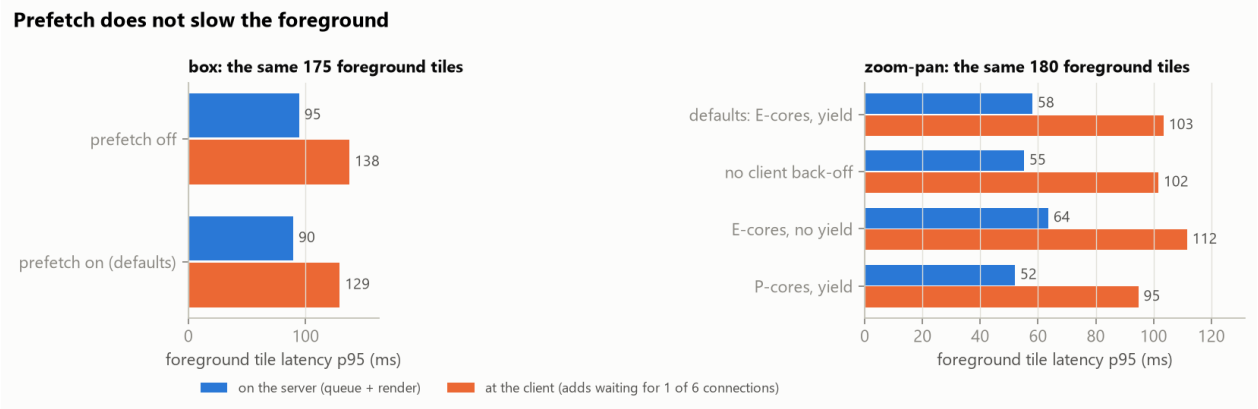


Figure 4: Foreground tile latency (95th percentile) on identical tile sets: on the server (queue + render) and at the client (also waiting for one of six connections).

The gate as specified compares foreground p95 per trace with prefetch on and off. That compares different tiles: with prefetch on, pans need almost no foreground tiles, and what remains are cold tiles at a new level. The like-for-like test is the box trace, where the foreground tiles are identical (Figure 4):

- **Result:** server p95 is 90 ms with prefetch and 95 ms without; client p95 is 129 vs 138 ms. Prefetch is slightly faster, because it has already decoded shared blocks. The gate passes.
- **Why:** prefetch stops at the first input event of a gesture and restarts only once the view is complete. Only 25 of 905 foreground tiles (3%) were ever requested while a prefetch request was in flight.
- **Placement:** letting prefetch start while foreground work is pending raises server p95 by 9% on zoom-pan.
- **Pinning held throughout:** 18,250 foreground jobs ran on performance cores only, and 114,962 prefetch jobs on efficiency cores only.

4.5 Workers, caches and the rest

- **The browser's six connections set the foreground latency.** The client sees about 40 ms more than the server at p95, because tiles wait for a connection. 4, 6 or 8 foreground workers give the same client latency (Table 5); 8 are kept for a second client such as QGIS.
- **The block cache matters:** without it the server decodes 6× more blocks and p95 rises 40%. 256 MiB suffices for these short traces.

Table 5: Mixed traces (pan-walk, zoom-pan, two bursts, box; 25 runs each; medians).

Configuration	Client p95 ms	Server p95 ms	Complete ms	Blocks decoded	Peak RSS MiB
defaults (8 workers, 1024 MiB block cache)	104	98	260	185	608
4 / 6 foreground workers	106 / 101	76 / 92	254 / 251	185	588 / 598
no block cache	146	111	285	1,128	386
256 MiB block cache	108	94	260	185	608
result cache off	103	91	257	185	576
slope / aspect	115 / 115	97 / 103	255 / 258	185	606 / 608
precomputed slope	93	56	248	160	562
HMA 8 m DEM	118	92	268	130	594
OS file cache flushed	103	104	257	185	608

- **The result cache** has no effect within a single fresh session, where no tile is asked for twice. It serves repeat visits and several clients.
- **Product, mode and data:** slope and aspect are about 10% slower than hillshade. Precomputed slope is fastest on the server (no halo, no kernel). The 8 m DEM behaves like the 2 m one.
- **A flushed OS cache changes nothing:** decoding, not the disk, is the cost, as in Phases 2–4.

4.6 A real browser

The same traces replayed in headless Chrome, driven in real time over the DevTools protocol, gave:

- the same commits, foreground requests (15, 0, 35) and prefetched tiles as the virtual client;
- view-complete times within 10% (298 vs 325 ms; 123 vs 127 ms);
- first frames 18–25 ms sooner in Chrome.

The simulated latencies are therefore slight upper bounds.

5 Discussion

Answers.

1. Yes: computing derivatives on demand keeps up with continuous viewing. Panning is seamless with a ring of neighbouring views; the first exact tile after a zoom arrives within about 90 ms of the commit.
2. The debounce removes two-thirds of the server's work per zoom gesture, for a fixed wait that 150 ms keeps short while still coalescing a notched wheel.
3. Prefetch costs efficiency-core time and 2–5 fetched tiles per tile used, but no foreground latency, because it never runs while the user is waiting.
4. The latency that remains comes from the browser's connection limit and from decoding compressed blocks, not from computing slope or hillshade. This is the same bottleneck Phases 2–4 found for whole-grid processing.

Limitations.

- One client on one laptop: the result cache and the spare workers serve several clients, which E6 does not run.
- Short, synthetic traces (1–11 s): caches never fill, so long sessions are not measured.
- Every run starts with an empty server, so a returning user would see faster tiles.

- CPU only: Phase 3 predicts no GPU gain at 512×512 per tile once the transfer is counted.

6 Next step

Phase 6 brings Phases 2–5 together in the final report, slides and a recorded demonstration of the viewer and the QGIS plugin.